

Study Plan and Progress

Claude Certified Architect - Foundations | exam Tue 25 August 2026

The desk companion. Two documents that change as you go: the calendar, and the record of what you have actually scored and where you keep going wrong. Reprint it whenever either changes - the markdown files are the source of truth, never this PDF.

Read section 2 before every practice set. The weak-point list is the whole value of it: each entry is a mark you have already lost at least once, written down so you do not lose it again. Weak point 19 is the one to carry into the exam.

Contents

1	The calendar
2	Progress log, scores and weak points

CCAR-F - 15-Day Study Plan

Exam: Claude Certified Architect - Foundations 60 questions - 120 minutes - you pass at 720 out of 1000

Exam date: Tuesday 25 August 2026. Study period: Saturday 8 August to Monday 24 August.

Book the exam now if you have not yet. Pearson VUE slots fill up, and you can cancel or move it free of charge up to 24 hours before. Booking early costs nothing and protects the date.

The daily shape - 2 hours

Use the same four blocks every day. The order matters: recall first, while you are fresh; practice last, when you are tired, because the exam is at the end of a long morning too.

Block	Time	What
1. Recall	10 min	Cheat sheet (one section) + Domain 4 facts. Cover the answers. Say them out loud
2. New material	50 min	Domain notes, or a course module
3. Practice	45 min	Drill questions, or /cert-exam
4. Review	15 min	Go through every wrong answer. Add one line to LOG.md

Block 4 is the one people skip. Do not. Reviewing a mistake teaches more than answering a new question correctly.

The calendar

Date	Day	Block 2 - new material	Block 3 - practice
Sat 8 Aug	[done]	Done - answer patterns	[done] Set 1: 11/20
Sun 9 Aug	-	Rest. Book the exam	-
Mon 10 Aug	1	Domain 1 notes 1.1-1.4 - start <i>Claude Code in Action</i>	/cert-exam 15 - Domain 1
Tue 11 Aug	2	Domain 1 notes 1.5-1.7 - <i>Claude Code in Action</i> (Agent SDK, hooks)	Domain 1 drill (15)
Wed 12 Aug	3	Domain 2 notes - start <i>Intro to MCP</i>	/cert-exam 15 - Domain 2
Thu 13 Aug	4	Finish <i>Intro to MCP</i>	Domain 2 drill (15)
Fri 14 Aug	5	[done] Domain 3 - Domain 4 and Domain 5 both pulled forward	[done] D4 ad-hoc 13/15 - D5 ad-hoc 14/15

Replanned 14 August - all five domains covered, 7 days early

The original calendar spent days 8-12 on first passes through Domains 4 and 5. Both are done, at 13/15 and 14/15. **Coverage is no longer the gap.** Two things are:

1. **Independence of source.** Every question answered well so far came from these notes, or was written from them. Purcell's set is the only measure that is not - keep it sealed until 23 Aug.

2. **Time pressure.** 60 questions in 120 minutes has never been rehearsed.

The freed days go to those two, and the drills stay where they are so they land as spaced retrieval rather than first passes.

Date	Day	Block 2 - new material	Block 3 - practice
Sat 15 Aug	6	Exercise 2 - configure a real repository (2 h, replaces blocks 2 and 3)	/cert-exam 15 - Domain 4 , then Domain 5 . Independent questions. Record both scores
Sun 16 Aug	7	-	Mock 1: /cert-exam 60, timed 120 min + full review. Cold. First honest cross-domain number
Mon 17 Aug	8	React to Mock 1 - re-read whatever it exposed, nothing else	/cert-exam 15 - the weakest domain in Mock 1
Tue 18 Aug	9	<i>Building with the Claude API</i> (tool use, prompt engineering)	Retake Domain 1 drill - the stalest domain, untouched since 10 Aug
Wed 19 Aug	10	<i>Building with the Claude API</i> (agents)	Domain 4 drill (15) - cold, 5 days after the ad-hoc set
Thu 20 Aug	11	Exercise 4	/cert-exam 15 - any domain not yet covered by a /cert-exam block
Fri 21 Aug	12	Re-read the two weakest domains from Mock 1	Domain 5 drill (15) - cold, 7 days after the ad-hoc set
Sat 22 Aug	13	Skim the six scenarios (20 min, no memorising)	Retake Domain 2 + 3 drills
Sun 23 Aug	14	-	Mock 2: Purcell's 60-question set, timed 120 min + full review
Mon 24 Aug	15	Read 03-cheat-sheet.md right through, twice	Redo only the questions you have ever got wrong
Tue 25 Aug	-	EXAM	-

Two things to hold on to through the last ten days. The daily Domain 4 facts do not stop just because Domain 4 now scores well - that drill is why it does. And **step 6 of the 30-second method** (ask where the information should have been produced, not where the failure showed up) is the single pattern behind every recent miss, across three domains.

What the extra time bought you

Compared with the 60-minute version:

- **Two timed mocks instead of one.** Mock 1 on 16 Aug is early enough to change the plan; Mock 2 on 23 Aug is the rehearsal.
- **Drill retakes at 7-day intervals.** Domain 1 on 11 and 18 Aug, Domains 2 and 3 on 13/15 and 22 Aug. Spaced repetition is what converts recognition into recall - and it is the method you already know from medical training.
- **The courses get real time**, instead of being squeezed between other work.

- A /cert-exam 15 on every domain the day you study it, on top of my drills. That is roughly 170 practice questions in total.

How the week is shaped

Weekdays are 2 hours. **Sat 15, Sun 16 and Sun 23 need 2-3 hours** - Exercise 2 and the two mocks do not fit in less.

What was cut, and what was protected

The plan started at 20 days. It is now 15. Domain 3 lost a day, because you already scored 4 out of 6 there and it is memorisation you drill daily anyway. The scenarios day was merged into the weak-domain day, and two review days were removed.

Domain 4 kept three days. It is 20% of the exam and you scored 1 out of 5. Do not shorten it, even if you fall behind.

Superseded 14 Aug. Domain 4 was covered in one sitting and scored 13/15, so those three days were released in the replan above. The daily fact drill is what made that possible and **still runs every day** - it is not what was cut.

If you fall behind

Drop things in this order: 1. The hands-on exercises (except Exercise 2) 2. The courses 3. Domain 3 study - the cheat sheet covers it

Never drop: the daily facts, the drills, or the Sunday 23 mock.

Practice questions

Source	Questions	Multiple-response?
Mine - Set 1 + five domain drills	95	Yes
Community /cert-exam bank	77	No
Purcell's practice set (1784098676646.pdf)	60	Yes - 11 of them

Roughly 230 practice questions. Purcell's set is the closest thing to the real exam: written by someone who sat and passed it, blueprint-weighted 16/11/12/12/9, and pitched at the difficulty he met on the day. Save it for 23 August and take it cold - reading it early destroys your only realistic rehearsal.

Time on exam day

60 questions in 120 minutes = **2 minutes per question**.

But each scenario has a story you read once and then reuse. So the real plan is:

- **Skim** each scenario story, do not study it - about 1 minute each = 4 minutes
- About **1 minute 55 seconds** per question = 116 minutes

The questions are self-contained. Any detail you need is repeated inside the question, so reading the story carefully is wasted time. A candidate who passed finished in **90 of the 120 minutes**.

About 1 question in 5 asks for two answers. In a 60-question exam that is roughly 11 of them. Read the line that tells you how many to select - it is the easiest mark to throw away.

If a question takes more than 3 minutes, mark it and come back later. There is no penalty for guessing, so never leave a question empty.

Exam facts to remember

- You get **4 scenarios**, chosen at random from a set of **6**. All 6 are in the guide.

- Questions are multiple-choice or multiple-response. **Each question tells you how many answers to select.** Read that line.
- Score is 100-1000. You pass at **720**. You are measured against a fixed standard, not against other people.
- Your score report shows percent correct per domain. But pass or fail depends only on the **total** score.
- The credential lasts **12 months**. Renewal is free if you do not let it expire.
- If you fail: wait 14 days, then 30 days, then 90 days. Maximum 4 attempts in 12 months.

Files in this folder

```

prep/
00-study-plan.md          <- this file
01-answer-patterns.md     <- read first. Re-read on day 19
02-course-map.md          <- which Partner Academy courses to take and skip
notes/
  domain-1-agentic-architecture.md    27%
  domain-2-tool-design-mcp.md         18%
  domain-3-claude-code-config.md      20%
  domain-4-prompt-structured-output.md 20% <- your weakest
  domain-5-context-reliability.md     15%
  scenarios.md                        all 6 exam scenarios
practice/
  set-01-questions.md      20 questions (day 1)
  set-01-answers.md        answers and explanations
  drill-domain-1.md        15, questions and answers in one file (days 2-4)
  drill-domain-2.md        15 (days 5-6)
  drill-domain-3.md        15 (days 7-9)
  drill-domain-4.md        15 - written 10 Aug, still cold. Take it 19 Aug
  drill-domain-5.md        15 - written 11 Aug, still cold. Take it 21 Aug
  daily-facts-domain-4.md   5 minutes, EVERY DAY from day 2
  set-02-*.md              30, mixed - written after your day-1 results
  set-03-*.md              60, full test - written after your day-10 results
exercises/
  README.md                the guide's 4 exercises, with checkpoints

```

Note on the two file types: the drills keep questions and answers in one file, so cover the answers while you work. The numbered sets keep them in two files, so you cannot see the answers during a timed test.

A note about language

The notes and explanations are written in simple English.

The **practice questions are not**. They use the same long, formal English as the real exam. This is on purpose - reading that style quickly is part of the test.

If a practice question is hard to understand, ask about it by number. Working out what a difficult question is really asking is useful practice for exam day.

Progress Log

Purpose: so any future session - with me or with a different assistant - can see where you are in about 30 seconds, without you having to explain again.

How to use it: after each study session, add a line to *Session history*. Update *Where I am now* and *Scores*. Keep it short. One or two sentences is enough.

Where I am now

Exam	
Grade	
Certified	
Architect	
-	
Foundations	
(CCAR-F)	
Exam	
date	
August	
2026	
Booked?	
booked	
11	
Aug,	
for	
Tue	
25	
August	
Study	
period	
August	
-	
Monday	
24	
August	
(15	
days)	

Current
day

of
15

-
Fri
14
Aug.

**All
five
domains
covered**

**and
above
target,**
seven
days
ahead
of
the
plan

Next
session

Aug
-

/cert-exam
15

on
D4

and
D5.

Independent
questions,
not

written
from

these
notes.

Then
Mock

1
on

16
Aug,

timed
and
cold

What attention
is across
still
untested
-
60
questions
in
120
minutes
-
questions
not
derived
from
these
notes.
See
weak
point
21

Study
time
hours
per
day

Scores

Date	Test	S	Notes
8 Aug	Set 1 (20, cold)	11/20	2/3 - D2 2/2 - D3 4/6 - D4 1/5 - D5 2/4

10 Aug	Domain 1 drill (15)	10/15 Missed Q7, Q9, Q13, Q14, Q15. Taken before the rebalance - see note below
10 Aug	D1 ad-hoc set (15)	10/15 ~ 14 of 15 answers were B. Discarded
10 Aug	D1 ad-hoc set, balanced (10)	10/10 ~ I missed both traps missed earlier: unknown-s decompos and resume-vs
11 Aug	Domain 2 drill (15)	14/15 Missed Q15 (Grep/Glob) Taken before the rebalance - 12 of 13 singles were B, so treat as ~12
11 Aug	D2 ad-hoc set, balanced (15)	13/15 Missed Q1 (error category), Q10 (tool granularity

11 Aug	D2 ad-hoc set, balanced (10)	7/10	Missed - read the questions too fast. Missed Q2, Q4 (half), Q10
11 Aug	D2 ad-hoc set, balanced (10)	9/10	Missed Q7 (tool_ch values). Error categories and search triad all clean
12 Aug	Mixed D1+D2 set (15, balanced)	12/15	Missed 8/9 - D2 4/6. Missed Q4 (tool granularity Q11 (decompos threshold), Q13 (semantic search). All three are repeat misses - see weak point 10
12 Aug	/cert-exam 15 - D2		

12 Aug	Domain 3 drill (15) - pulled forward from 15 Aug	13/15 Missed Q2 (@import - 2nd time) and Q12 (gave one answer where two were asked). Memorisation core clean: all paths, frontmatter keys and flags correct
14 Aug	/cert-exam 15 - D3	
14 Aug	D5 ad-hoc set, balanced (15)	14/15 Missed Q4 only (subagent metadata - chose <i>both</i> wrong options). Both escalation directions clean in one sitting

14 Aug	D4 ad-hoc set, balanced (15)	13/15 Used Q6 (required-v reverse direction) and Q11 (format normalisat in the prompt). Both are Block A schema facts. All three select-two clean
--------	------------------------------	---

14 Aug	Mixed D1+D2+D3 set (20, balanced)	17/20 6/8 - D2 6/6 - D3 5/6. Missed Q1 (stopping the loop on "text and no tool_use"), Q5 (one answer where two were asked - 5th time), Q15 (@import - 3rd time , same wrong answer as 12 Aug)
~~15 Aug~~	~~Domain 3 drill (15)~~ - taken 12 Aug	Retake on 22 Aug as planned
16 Aug	Mock 1 - /cert-exam 60, timed	Record the score for every domain
17 Aug	/cert-exam 15 - D4	

18 Aug	Domain 1 drill - retake		Compare with 10 Aug (10/15, pre-rebala First honest score on these questions
19 Aug	Domain 4 drill (15)		
20 Aug	/cert-exam 15 - D5		
21 Aug	Domain 5 drill (15)		
22 Aug	Domain 2 + 3 drills - retake		
23 Aug	Mock 2 - Purcell's 60, timed		The rehearsal. Take it cold

Target before the exam: 80%+ on practice sets, and no domain below 60%.

On the retakes: the score matters less than *which* questions you miss twice. A question missed twice is a fact you have not learned. Put it on the cheat sheet.

What I know about my own weak points

Updated as evidence appears. Right now, from Set 1 only:

~~**Domain 4 is the big gap - 1 out of 5, on 20% of the exam.**~~ **Largely closed 14 Aug (13/15).** Kept here because the diagnosis below still explains what to keep doing - see weak points 16 and 17 for what remains. Three of the four misses were flat facts: `tool_choice` values (auto allows text), required schema fields causing invented values, and which errors a retry can fix. -> Fixed by `daily-facts-domain-4.md`, 5 minutes every day. The fourth (Q17, voting across passes) was applying a rule too widely - see point 2. So the Domain 4 gap is not pure recall.

Applying a rule too widely. Chose plan mode for a single-file fix (Q12). Chose voting across review passes (Q17), which is on the "never correct" list. -> Re-read `01-answer-patterns.md`.

Mechanism confusion. Mixed up `fork_session` with parallel Task calls (Q8). Blamed the wrong component when the question named it (Q5).

The escalation trap (Q3). A customer who explicitly asks for a human = escalate immediately. A customer who is frustrated but can be helped = offer help first. These two look similar and the exam uses that.

What is already working: structured error design, programmatic enforcement vs prompting, two-layer context management, conflicting sources, batch vs real-time. The architecture reasoning is sound. The gap is recall.

Added 10-11 Aug, from the Domain 1 and 2 drills

Rules stored where the guide has a threshold. Three Domain 1 misses were all this shape: *stale* -> *fresh session* (it is actually "how much is stale"), *audit* -> *per-file passes* (it is actually "do I know the shape before I start"), and `!= end_turn` -> *continue* (it is actually "continue only on `tool_use`"). Each rule works most of the time, which is why it sticks. -> Convert each into a question you ask the scenario. All three came out clean afterwards.

Stopping at the first correct answer on select-two questions. Happened three times. -> Reliable tell: **if the scenario names two distinct symptoms, the two answers map onto them.** One answer cannot be complete.

Error categories decided by feel rather than by test. Two misses. The test, in order: did the system fail and might a retry succeed -> *transient*; is the input malformed -> *validation*; input and system both fine but a rule says no -> *business*; caller lacks access -> *permission*. Only transient is retrievable.

Grep / Glob / semantic search. Missed three times before it stuck. **Glob = file paths. Grep = text you can spell exactly. Semantic = concepts with no literal to search for.** If the question prints the string you are looking for, it is Grep.

tool_choice values. auto = may call a tool *or answer directly*. any = must call some tool, model picks which. {"type": "tool", "name": ...} = must call that one. "any" reads like "anything permitted"; it means "one of them, mandatory."

Confirmed strong in Domain 2: programmatic enforcement, scope partitioning vs iterative refinement, MCP resources vs tools, .mcp.json vs ~/.claude.json, \${VAR} expansion, community vs custom servers, local recovery vs reporting upward.

Added 12 Aug, from the mixed D1+D2 set (12/15)

Three facts are now missed twice or more. These are the cheat-sheet items.

Fact	Missed	Now on the cheat sheet as
Chaining vs dynamic decomposition	3x (10 Aug, and Q11 here)	"Do I know the shape before I start?" - with the note that <i>audit / large / comprehensive do not decide it</i>
Grep / Glob / semantic	4x	"If the question says you do not know what it is called -> semantic"
Tool granularity (<code>mode</code> enum -> split)	2x	"Merging the modes into fewer modes is the wrong answer"

On Q11 the material was partly at fault: the cheat sheet had **no** chaining-vs-dynamic entry at all, and its one line on the subject ("large review -> one pass per file") sat in the Domain 4 section with no qualifier, priming exactly the wrong answer. Fixed 12 Aug in both the cheat sheet and the revision card.

Evidence the rule itself is known: Q12 (same rule, known 22-file list) was correct in the same sitting. What fails is the discriminator, not the rule.

What is working: all three select-two questions clean. Weak point 5's resume half is closed too (Q10 correct, the exact trap that cost a mark on 10 Aug). `fork_session` vs `parallel` Task calls clean (Q6), so weak point 3 is closing.

Added 12 Aug, from the Domain 3 drill (13/15)

Weak point 6 is improved but NOT closed - this is now its 4th occurrence. Domain 3 Q12 asked for two answers and got one. Three of the four multi-answer questions in that drill were right, and all three were right in the mixed set earlier the same day, so the habit is mostly there. But the failure mode has not changed shape once. -> **Exam-day rule: before answering any question, look at whether it says "select two".** If it does, count the symptoms in the scenario, then check you have one answer per symptom. Do this as a mechanical step, not as a judgement.

@import is a positive rule that has not anchored - missed twice. Set 1 Q9 and Domain 3 Q2 are the same decision. The improvement: the old wrong answer (path globs) was correctly rejected this time; a new wrong answer (skills) replaced it. -> **@import = the maintainer selects which shared documents apply, with no duplication.** Skills must be **invoked**, so they can never hold standards that must always apply. Path globs are for conventions that follow **file type across many folders**.

Added 14 Aug, from the mixed D1+D2+D3 set (17/20)

@import is now missed three times, twice with the same wrong answer (skills). This is the most reliable predictor of a lost mark in the whole repo. The pattern of the miss has stabilised: the wrong answer is always *skills*, chosen because

the question describes a person **choosing** which documents apply, and "choosing" feels like "invoking". -> **The discriminator is WHO chooses and WHEN.** @import: the *maintainer* chooses once, when writing the file, and the content then loads **always**. A skill: the *model or developer* chooses at run time, every time, and the content loads **only when invoked**. If the question says the standards must apply to all work in that package, a skill cannot do it - no matter who selected it.

A new shape of weak point 5: substituting an inference for the actual signal. Q1 chose "exit when the response has a text block and no `tool_use` blocks". That condition is *usually true* at the end of a turn, which is exactly why it is attractive - same failure shape as the three misses in weak point 5. But the guide names "checking whether the response contains text" as one of the four wrong ways to stop. -> **The loop reads stop_reason and nothing else.** Inspecting block types is a proxy, and the exam offers proxies as wrong answers on purpose. Note the second half of the trap: the replacement in the question was *also* wrong (`task_complete`), so the question was testing whether you would swap one wrong mechanism for another.

Weak point 6 is at its 5th occurrence - but the failure has moved. Q5 was answered with one letter; Q10, Q14 and Q20 were all answered with two, correctly. So the counting habit works when the two symptoms are *the same kind of thing*. Q5's two symptoms were different kinds - a policy breach and a data-format problem - and the second one was found while the first was dropped. -> **Extend the exam-day rule:** after finding one answer, do not ask "is this right?" but "which symptom did this answer solve, and what solves the other one?"

Added 14 Aug, from the D4 ad-hoc set (13/15)

Domain 4 is no longer the weakest domain - 1/5 cold on 8 Aug, 13/15 now. Everything the daily fact drill targets came out clean: `tool_choice`, the batch table (including no multi-turn tool calling), what a retry request contains, few-shot count and targeting, and both confidence traps in the same sitting - confidence **next to** a finding for triage is correct, confidence as a **filter** on what gets reported is not. -> Keep the daily drill running; do not declare this closed until the 16 Aug mock agrees.

Both misses were Block A schema design, and both were the same error: taking a rule in one direction only.

- *Required vs nullable.* The learned rule is "required makes the model invent values -> use nullable". The scenario was the reverse case - a machine-generated document that **always** contains the field, where the failure was the field being **omitted**. The answer is **required**. The chosen answer, "nullable plus a validation rule rejecting null", gets the failure mode of required with none of the guarantee. -> **Discriminator: does the source always contain this value?** Yes -> required. Might not -> nullable. Both are correct answers on this exam. Added to the cheat sheet and revision card.

- *Format normalisation.* Chose a post-processing layer over normalisation rules in the prompt. Defensible engineering, wrong exam answer. -> **The schema fixes the type, the prompt fixes the format.** The cheat sheet was missing this line entirely - the revision card had it, the cheat sheet did not. Fixed. Same class of material gap as the one Q11 exposed on 12 Aug, so re-check the two files against each other rather than assuming they match.

Note what this is **not**: Q3 ("unclear") and Q9 (`conflict_detected` plus extracting both readings) came from the same block and were right. It is two specific rows, not the block.

Weak point 6 is clean twice running. All three select-two questions correct here, after 3 of 4 on 14 Aug's mixed set. The mechanical "count the symptoms" step is working, including on Q14, where the two symptoms were **different kinds** - the shape that failed on 14 Aug.

Added 14 Aug, from the D5 ad-hoc set (14/15) - and this is the one to carry to the exam

All five domains are covered, and one error shape now accounts for every recent miss: repairing a symptom downstream instead of requiring the right structure upstream. Three misses across 30 questions, in three different domains, all the same:

Domain	The failure	Chose (late fix)	Correct (early fix)
D4	A field the source always has goes missing	nullable + reject null	required
D4	Dates and currency arrive inconsistent	a post-processing layer	normalisation rules in the prompt

D5	Synthesis cannot judge a source's age or type	re-fetch at synthesis; weigh credibility	subagents record dates, location, methodology
----	---	--	--

Every late fix is **defensible engineering**, which is exactly why it is offered. The exam is testing whether you locate the component that held the information at the moment it was lost. -> Now step 6 of the 30-second method, in the cheat sheet and on the revision card.

Note the failure mode on D5 Q4: both selected answers were wrong, and both mapped cleanly onto the two symptoms in the scenario. So the select-two counting habit worked and produced two answers from the wrong family. Counting symptoms and judging content are separate steps; the first being right says nothing about the second.

Domain 5 is covered: 14/15, with both escalation directions correct in one sitting. No progress possible + calm customer -> escalate (Q2). Explicit request for a human mid investigation -> escalate **immediately**, do not finish the diagnosis first (Q13). Criteria plus few-shot rather than a classifier (Q8). The Set 1 escalation trap is closed. Q15 was an over-application trap - a vague answer in a *fresh* session, where nothing has degraded and the agent simply never read the file - and it was not taken, which is weak point 2 improving.

Read the two 14 Aug scores with their caveat. Both sets were written from these notes and taken in the same sitting as the review, so they measure recall at zero delay from material just read. The cold numbers are still Set 1 (11/20) and whatever Mock 1 returns. More generally: every question answered well so far traces back to one source - these notes, from guide v1.0, and questions written from these notes. **Purcell's set is the only independent measure, which is the strongest argument for keeping it sealed until 23 Aug.**

Session history

Date	What happened
14 Aug	Domain 5 straight after Domain 4: 15 new questions, balanced key, all six task statements, avoiding <code>drill-domain-5.md</code> so it stays cold for 21 Aug. 14/15 - one miss. All five domains are now covered, seven days ahead of the plan. Reworked the calendar from 15 Aug onward: the four days reserved for domains already finished go to independent question sources and timed practice instead, because coverage is no longer the gap - independence of source and time pressure are. Added weak points 19-21; 19 is the one to carry to the exam and is now step 6 of the 30-second method.
14 Aug	Started Domain 4, the weakest domain. Ad-hoc set of 15 new questions, balanced key, all six task statements, deliberately avoiding the framings in <code>drill-domain-4.md</code> so that drill stays cold for 19 Aug. 13/15 - up from 1/5 cold on 8 Aug, and above the drill's own "done with this domain" bar. Both misses were Block A schema design, both the same shape: applying a rule in one direction only (weak point 17). Fixed the cheat sheet, which stated "make it optional and nullable" with no qualifier and had no format-normalisation line at all; the revision card had the second but not the first. Both regenerated.

14 Aug	Mixed set across all three completed domains: 20 new questions, blueprint-weighted, balanced key. 17/20 (D1 6/8, D2 6/6, D3 5/6) - clears the 80% target with every domain above the floor, and Domain 2 is now clean twice running. All three misses were traps aimed at the repeat-miss list: <code>@import</code> (3rd time, same wrong answer as 12 Aug), a select-two answered with one (5th time, but 3 of 4 select-twos were right), and a <i>new</i> variant of weak point 5 - stopping the loop on "text and no <code>tool_use</code> " instead of on <code>stop_reason</code> . Weak points 13-15 added. Notable: Grep-vs-semantic and tool granularity, both long-running repeat misses, were correct when asked from the reverse direction.
12 Aug	Domain 3 drill, pulled forward from 15 Aug: 13/15 , clearing the drill's own "done with this domain" bar. The memorisation core - file paths, frontmatter keys, CLI flags - was fully clean, which is the bulk of this domain. Two misses, both recurrences: <code>@import</code> (2nd time) and a select-two answered with one answer (4th time). Added two things to the cheat sheet in response: a "which mechanism holds which content" table, and a mechanical select-two check.
12 Aug	Confirmation set before starting Domain 3: 15 mixed D1+D2 questions, balanced key, all new. 12/15 (D1 8/9, D2 4/6) - clears the 80% target with both domains above the 60% floor. The three misses were all repeat misses, so they went onto the cheat sheet; one of them (decomposition) exposed a real gap in the cheat sheet itself, now fixed. Select-two habit confirmed closed. Cleared to start Domain 3.
11 Aug	Booked the exam for Tue 25 August. Built the Domain 5 drill (15 questions, all six task statements, balanced key) - every domain now has a drill.
11 Aug	Cross-checked Purcell's 60-question set and the official exam guide (both PDFs were already in the folder). Purcell's answer key agrees with this repo's material on all 60 questions - zero contradictions, though topical overlap is heavy since both are written against guide v1.0. The guide adjudicated every remaining docs-drift item in <code>AUDIT.md</code> in the notes' favour; annotations added, <code>stop_reason</code> count corrected to seven, <code>AUDIT.md</code> is now fully closed. Bonus finding: the guide itself uses both <code>isRetryable</code> and <code>retriable</code> in the same task statement - the origin of the spelling inconsistency.
11 Aug	A second review independently confirmed every verified <code>AUDIT.md</code> finding, then applied the fixes: the <code>--json</code> -schema contradiction, one spelling for the retriable field (<code>isRetryable</code>), the Set 1 blueprint header, the Domain 4 miss count (four, not three - in four files, including this one), the resume-vs-fresh "how much is stale" qualifier (now also on the cheat sheet and revision card), required-vs-nullable, Scenario 4's missing Domain 5, the fourth wrong-way-to-stop (<code>task_complete</code>), and the <code>PreToolUse</code> name. All three PDFs regenerated. Docs-only items 3a-3e/3g/3h left for a session with the exam guide open.

11 Aug	Domain 2: took the repo drill (14/15) then three balanced ad-hoc sets (13/15, 7/10, 9/10). Discovered a severe answer-key bias in all five practice files - 56 of 66 single-answer questions were B and D was never correct, so "always answer B" scored 56/66. Rebalanced every file and verified no option text changed. Ran a full audit of the material for contradictions and factual drift -> AUDIT.md , nothing fixed yet.
10 Aug	Domain 1: took the repo drill cold (10/15), reviewed the whole domain by mental model rather than fact list, then 10/10 on a balanced ad-hoc set. Three weak spots identified and closed - see weak points 5-7 above.
10 Aug	Reviewed Matthew Purcell's LinkedIn exam review and his 60-question practice set. Confirmed three earlier calls (Bedrock/Vertex out of scope, Claude 101 and AI Fluency too basic, exam is memorisation-heavy). Downgraded the scenarios from "study" to "skim once". Adjusted exam-day timing. His set replaces the mock I was going to write for 23 Aug.
8 Aug	Identified two different exams (CCAR-F vs CCAR-P) and chose Foundations. Confirmed Foundations is accepted for Intertec partner tier. Built the full prep folder. Took Set 1 cold: 11/20. Converted everything to simple English. Found and installed the community 77-question practice exam as /cert-exam. Pushed everything to GitHub.

Still to be built

Needed by	What
~~19 Aug~~	~~Domain 4 drill~~ - built 10 Aug
~~21 Aug~~	~~Domain 5 drill (15 questions)~~ - built 11 Aug . Balanced key, covers all six task statements, deliberately avoids Purcell's D5 framings so the 23 Aug mock stays cold
~~Before revision day~~	~~Work through AUDIT.md~~ - fully closed 11 Aug . Verified items fixed, docs-drift items adjudicated against the exam guide (the notes matched the guide on every row), annotations added, PDFs regenerated

~~23 Aug: full mock~~ - **no longer needed**. Matthew Purcell's 60-question practice set (1784098676646.pdf, kept locally, not in the public repo) replaces it. Written by someone who passed CCAR-F, blueprint-weighted, and it includes 11 multiple-response questions.

External resources

Resource	What it is	When
/cert-exam	77 community questions by Olivier Legris, installed as a Claude Code skill. Single-answer only	Daily, 15 at a time

1784098676646.pdf	Matthew Purcell's 60-question set. Blueprint-weighted, 11 multiple-response, answer key with rationales. Cross-checked 11 Aug: zero contradictions with the repo material. Caveat: heavy topical overlap with the repo drills (same guide, same scenarios), so a high mock score partly reflects repetition - weight the domains you haven't drilled recently	23 Aug only - take it cold
Matthew Purcell's LinkedIn review	He passed all four exams. Confirms: Bedrock/Vertex not tested, Claude 101 and AI Fluency too basic, CCAR-F is the hardest of the four, scenarios are "dressing", finished in 90 of 120 min	Read once
Paul Larionov's GitHub repo	A study guide Purcell recommends. Not yet evaluated	Optional

To check: Purcell reports a 100% exam discount for Global Premier partners until 31 Aug 2026. The exam is 25 Aug, inside that window. Ask Intertec whether it is free.

Decisions already made - do not revisit

- **Foundations (CCAR-F), not Professional (CCAR-P).** Confirmed as accepted for partner tier.
- **Skip the Bedrock and Google Cloud courses.** Cloud provider configuration appears exactly once in the 39-page guide, in the out-of-scope list.
- **Skip AI Fluency and Claude 101.** Level 100, below this exam's candidate profile.
- **Take:** Claude Code in Action (covers Claude Code *and* the Agent SDK), Intro to MCP, and only the tool use / prompt engineering / agents modules of Building with the Claude API.
- **The official exam guide PDF is not pushed to GitHub.** It is Anthropic partner material and the repository is public. It is in `.gitignore`. The same applies to Purcell's practice set - it is his copyrighted work.
- **The scenarios are context, not content to memorise.** A candidate who passed calls them "more dressing than anything substantive". Skim them once; do not study them.

Useful context for a new session

- Write in plain English. English is not my first language. Keep exact exam terminology.
- Explanations in simple English; **practice questions stay at exam-level English on purpose**, because reading that style quickly is part of the test.
- Background: medical doctor, now in product and program management. Comfortable with architecture reasoning and trade-offs; less comfortable with backend implementation. This exam does not test backend implementation - Section 17 of the guide excludes it.
- Files are mirrored at github.com/abracadaniel92/claude_architect_exam.